The background of the slide features a close-up, shallow depth-of-field photograph of a molecular model. The model consists of dark blue spheres connected by thin, coiled metal springs, representing atoms and bonds. In the background, a portion of a periodic table is visible, showing elements like Cu, Zn, Ag, Cd, Au, and Hg. A dark, semi-transparent horizontal band is positioned across the lower half of the image, serving as a backdrop for the text.

## **AMS 530: Final Project 4.4**

### **Molecular Dynamics Energy Reduction**

Amr Ibrahim

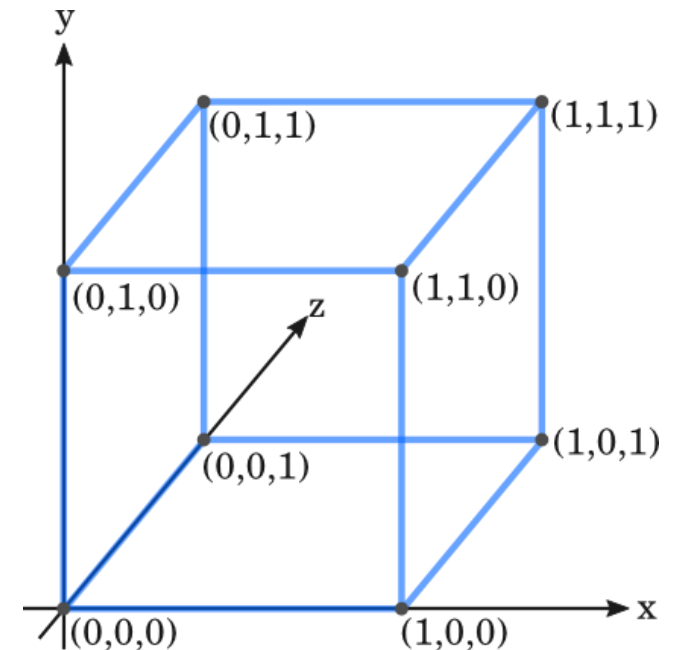
# Table of contents

- Problem statement
  - Goal
  - Setup
- Program
  - General Design
  - Functions spotlight
- Results
- Discussion
- Future direction



# Problem Statement

- Goal: Reduce energy of given system of particles by kicking them around
  - Emphasis on parallelization of energy reduction
- Total Energy is given by
  - $\frac{1}{2} \sum_{j \neq i}^N V_{ij}$
  - Lenard-Jones potential (ignoring constants)
    - $V_{ij} = \frac{1}{r_{ij}^{12}} - \frac{2}{r_{ij}^6}$
    - $r_{ij} = |x_i - x_j|$
- Initial conditions:
  - Box 10x10x10
  - $N = 11 \times 11 \times 11$  particles placed at box grid points



# Program: General Design

1. Generate initial coordinates & initial energy  $E_0$ 
  - Initial coordinates generated by all processes
  - Particles distributed among processes
  - Energy calculation
2. Move a particle randomly with max distance 0.5
  - Moving a fraction of particles at a time
  - Discard moves that do not reduce energy
  - If particle moves outside the box return it from the opposite side
3. Repeat step 2 for 100 steps and print out energy if lowered

# Program

- Distribute particles
  - Based on  $N // \text{size}$
  - Add remainder particles evenly across processes
- Generate initial coordinates
  - For all processes
  - All processes require full coordinates to compute initial energy
- Compute local energy
  - Compute interactions of local particles
  - Total energy =  $\sum_{i=0}^P \text{local energy}_i$

```
def distribute_particles(rank, size, N):  
    # Calculate the number of particles for each process  
    particles_per_process = N // size  
    remainder_particles = N % size  
  
    # Calculate the start and end indices for each process  
    start_index = rank * particles_per_process + min(rank, remainder_particles)  
    end_index = start_index + particles_per_process + (rank < remainder_particles)  
  
    return start_index, end_index
```

```
def generate_initial_coordinates(box_side):  
    grid_points = box_side + 1  
  
    # Generate grid points for each dimension  
    x, y, z = np.meshgrid(np.linspace(0, box_side, grid_points),  
                          np.linspace(0, box_side, grid_points),  
                          np.linspace(0, box_side, grid_points),  
                          indexing='ij')  
  
    # Flatten the grid points  
    coordinates = np.column_stack((x.flatten(), y.flatten(), z.flatten()))  
    return coordinates
```

```
def compute_local_energy(start_index, end_index, coordinates):  
    N = len(coordinates)  
    local_energy = 0.0  
  
    for i in range(start_index, end_index):  
        for j in range(N):  
            if i != j:  
                r_ij = np.linalg.norm(coordinates[i] - coordinates[j])  
                local_energy += 0.5 * (1 / (r_ij**12 + 1e-6) - 2 / (r_ij**6 + 1e-6))  
  
    return local_energy
```

# Program

```
def move_local_particles(start_index, end_index, coordinates, move_size, box_side, fraction):  
    moved_coordinates = coordinates.copy()  
  
    # Calculate the number of particles to move  
    num_particles_to_move = int((end_index - start_index) * fraction)  
  
    # Select particles to move  
    particles_to_move = np.random.choice(np.arange(start_index, end_index), size=num_particles_to_move, replace=False)  
    moved_coordinates[particles_to_move, :] += np.random.uniform(-move_size, move_size, size=(num_particles_to_move, 3))  
    moved_coordinates = np.mod(moved_coordinates, box_side) # Ensure particles stay within the 10x10x10 box  
    return moved_coordinates
```

- Move local particles
  - Move a fraction of randomly chosen particles
  - Move with distribution  $U(-0.5, 0.5)$
  - Mod coordinates
- Metropolis criterion (Simulated Annealing)
  - Update local coordinates & local energy if
    - $E_{local_i} - E_{local_0} < 0$
    - $rand < e^{-(E_{local_i} - E_{local_0})/T}, T = 1$
  - In case stuck at local min

```
def metropolis_criterion(delta_energy, temperature=1.0):  
    return delta_energy < 0 or np.random.rand() < np.exp(-delta_energy / temperature)
```



# Program

- Reduce energy
  - Move particles
  - Compute local energy
  - Metropolis criterion
  - Extract new local coordinates
    - Each process get their respective local particle coordinates
  - All gather new coordinates
    - Updated coordinates for next iteration
  - Compute new total energy
    - New total energy =  $\sum_{i=0}^P local\ energy_i$

```
def reduce_energy(comm, rank, start_index, end_index, initial_local_energy, initial_coordinates,
                  move_size, box_side, max_iterations):
    current_coordinates = initial_coordinates.copy()
    current_local_energy = initial_local_energy.copy()
    last_printed_energy = initial_local_energy.copy()

    iteration = 0
    while iteration < max_iterations:
        local_coordinates = move_local_particles(start_index, end_index,
                                                current_coordinates, move_size, box_side, fraction=0.5)
        local_energy = compute_local_energy(start_index, end_index, local_coordinates)

        if metropolis_criterion(local_energy - current_local_energy):
            current_coordinates = local_coordinates
            current_local_energy = local_energy

        sub_local_coordinates = local_coordinates[start_index:end_index]
        all_coordinates = comm.allgather(sub_local_coordinates)
        current_coordinates = np.concatenate(all_coordinates)

        current_energy = comm.reduce(current_local_energy, op=MPI.SUM, root=0)
        if rank == 0 and current_energy < last_printed_energy:
            print("Iteration {}: Energy = {}".format((iteration+1), current_energy))
            last_printed_energy = current_energy

        iteration += 1

    return current_coordinates, current_local_energy
```

# Results

P = 4  
Fraction = 0.1  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -360901528.832  
Final Energy: -360901528.832  
Time elapsed: 171.216108799

Fraction = 0.5  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -140645019.075  
Final Energy: -140645019.075  
Time elapsed: 171.37024641

Fraction = 0.9  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -6010.79030306  
Final Energy: -6010.79030306  
Time elapsed: 171.45086503

P = 16  
Fraction = 0.1  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -388855023.323  
Final Energy: -388855023.323  
Time elapsed: 46.7822186947

Fraction = 0.5  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -171426352.371  
Final Energy: -171426352.371  
Time elapsed: 46.2184393406

Fraction = 0.9  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -8491270.63421  
Final Energy: -8491270.63421  
Time elapsed: 46.3413505554

P = 64  
Fraction = 0.1  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -383651128.644  
Iteration 2: Energy = -383651841.792  
Iteration 6: Energy = -383925203.181  
Final Energy: -383925203.181  
Time elapsed: 13.6498684883

Fraction = 0.5  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -198968968.446  
Iteration 2: Energy = -199445840.306  
Iteration 31: Energy = -199561757.549  
Final Energy: -199561757.549  
Time elapsed: 13.0203437805

Fraction = 0.9  
Initial Energy: -6010.79030306  
Iteration 1: Energy = -22571662.3707  
Iteration 14: Energy = -23023038.0875  
Iteration 30: Energy = -23070678.4131  
Iteration 41: Energy = -23233142.6831  
Iteration 53: Energy = -23246182.9608  
Iteration 58: Energy = -23429096.1015  
Iteration 63: Energy = -23455705.3421  
Iteration 71: Energy = -23570977.2614  
Iteration 74: Energy = -24262885.7561  
Iteration 79: Energy = -24271188.6698  
Iteration 86: Energy = -24657573.0983  
Iteration 99: Energy = -24675107.7461  
Final Energy: -24675107.7461  
Time elapsed: 13.0429084301



# Discussion

- Energy reduction
  - 1<sup>st</sup> step usually has a massive reduction
  - As fraction number increases reduction of energy is reduced
  - Among iterations reduction is small compared to 1<sup>st</sup> iteration
- Time elapsed
  - Within each P, time elapsed variance is very small
  - As P increases time elapsed decreases

```
P = 64
Fraction = 0.1
Initial Energy: -6010.79030306
Iteration 1: Energy = -383651128.644
Iteration 2: Energy = -383651841.792
Iteration 6: Energy = -383925203.181
Final Energy: -383925203.181
Time elapsed: 13.6498684883

Fraction = 0.5
Initial Energy: -6010.79030306
Iteration 1: Energy = -198968968.446
Iteration 2: Energy = -199445840.306
Iteration 31: Energy = -199561757.549
Final Energy: -199561757.549
Time elapsed: 13.0203437805

Fraction = 0.9
Initial Energy: -6010.79030306
Iteration 1: Energy = -22571662.3707
Iteration 14: Energy = -23023038.0875
Iteration 30: Energy = -23070678.4131
Iteration 41: Energy = -23233142.6831
Iteration 53: Energy = -23246182.9608
Iteration 58: Energy = -23429096.1015
Iteration 63: Energy = -23455705.3421
Iteration 71: Energy = -23570977.2614
Iteration 74: Energy = -24262885.7561
Iteration 79: Energy = -24271188.6698
Iteration 86: Energy = -24657573.0983
Iteration 99: Energy = -24675107.7461
Final Energy: -24675107.7461
Time elapsed: 13.0429084301
```

# Future Direction

- Decomposition
  - This project uses particle decomposition
  - Spatial + particle decomposition
- Energy reduction
  - Method applied Monte Carlo
  - Gradient descent
  - Python's `scipy.optimize()`

